

1 Hands On: Association Rules

1.1 Association Rules

1. Load the packages `arules`, `arulesViz` and the dataset `Groceries` from the package `arules` which contains 1 month of real-world point-of-sale transaction data from a typical local grocery.

```
library(arules)
library(arulesViz)
library(dplyr)
data(Groceries)
```

- (a) Type `Groceries` on the R prompt. What does it return? Use the function `class` to inspect the type of data set.

```
Groceries

## transactions in sparse format with
## 9835 transactions (rows) and
## 169 items (columns)

class(Groceries)

## [1] "transactions"
## attr(,"package")
## [1] "arules"
```

- (b) Use the function `summary` to get more information on the data set.

```
summary(Groceries)

## transactions as itemMatrix in sparse format with
## 9835 rows (elements/itemsets/transactions) and
## 169 columns (items) and a density of 0.02609146
##
## most frequent items:
##      whole milk other vegetables      rolls/buns      soda
##           2513           1903           1809           1715
##      yogurt      (Other)
##           1372           34055
##
## element (itemset/transaction) length distribution:
## sizes
##      1      2      3      4      5      6      7      8      9     10     11     12     13     14     15
## 2159 1643 1299 1005  855  645  545  438  350  246  182  117  78  77  55
##      16      17      18      19      20      21      22      23      24      26      27      28      29      32
##      46      29      14      14      9      11      4      6      1      1      1      3      1
##
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.000  2.000  3.000  4.409  6.000 32.000
##
## includes extended item information - examples:
##      labels level2      level1
## 1 frankfurter sausage meat and sausage
## 2  sausage sausage meat and sausage
## 3  liver loaf sausage meat and sausage
```

- (c) Use the function `size` on the data set. What information does it return?

```
head(size(Groceries))  
  
## [1] 4 3 1 4 4 5
```

- (d) Use the function `inspect` to see the first five transactions.

```
inspect(Groceries[1:5])  
  
##      items  
## [1] {citrus fruit,  
##      semi-finished bread,  
##      margarine,  
##      ready soups}  
## [2] {tropical fruit,  
##      yogurt,  
##      coffee}  
## [3] {whole milk}  
## [4] {pip fruit,  
##      yogurt,  
##      cream cheese ,  
##      meat spreads}  
## [5] {other vegetables,  
##      whole milk,  
##      condensed milk,  
##      long life bakery product}
```

- (e) Are there any duplicated transactions? Use the function `unique` or `duplicated`.


```
length(which(duplicated(Groceries)))  
  
## [1] 2824
```

- (f) Use the function `itemFrequency` to see the relative frequency of each item.

```
head(itemFrequency(Groceries))  
  
##      frankfurter      sausage      liver loaf      ham  
##      0.058973055      0.093950178      0.005083884      0.026029487  
##      meat finished products  
##      0.025826131      0.006507372
```


- (g) Using the function `itemFrequencyPlot`, plot the top 5 more frequent items.

```
itemFrequencyPlot(Groceries, topN = 5)
```



- (h) Using the same function `itemFrequencyPlot`, plot the items that have a support value of at least 0.1. How many are there?

```
itemFrequencyPlot(Groceries, support = 0.1)
```



- (i) Using function `apriori`, and without generating any rules, obtain the frequent itemsets for a minimum support of 0.01. What is the class of the object returned? How many frequent itemsets were found?

```
fsets <- apriori(Groceries, parameter = list(supp = 0.01, target = "frequent itemsets"))

## Apriori
```

```
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          NA      0.1      1 none FALSE          TRUE      5      0.01      1
## maxlen      target      ext
##      10 frequent itemsets FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE  FALSE TRUE      2      TRUE
##
## Absolute minimum support count: 98
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.01s].
## sorting and recoding items ... [88 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 done [0.00s].
## writing ... [333 set(s)] done [0.00s].
## creating S4 object ... done [0.00s].

class(fsets)

## [1] "itemsets"
## attr("package")
## [1] "arules"
```

- (j) Inspect the 5 most frequent itemsets. What's their size?

```
inspect(sort(fsets)[1:5])

##      items      support      count
## [1] {whole milk}      0.2555160 2513
## [2] {other vegetables} 0.1934926 1903
## [3] {rolls/buns}       0.1839349 1809
## [4] {soda}             0.1743772 1715
## [5] {yogurt}           0.1395018 1372
```

- (k) From the frequent itemsets obtained, select the subset of closed frequent itemsets and the subset of maximal frequent itemsets. What can you conclude?

```
fsets[is.closed(fsets)]

## set of 333 itemsets

fsets[is.maximal(fsets)]

## set of 243 itemsets
```

- (l) Use the function `apriori` to generate association rules from the Groceries data set. What is the class of the returned object? How many rules were generated?

```
rules <- apriori(Groceries)

## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          0.8      0.1      1 none FALSE          TRUE      5      0.1      1
## maxlen target      ext
##      10 rules FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE  FALSE TRUE      2      TRUE
##
## Absolute minimum support count: 983
##
## set item appearances ...[0 item(s)] done [0.00s].
```

```
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.00s].
## sorting and recoding items ... [8 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 done [0.00s].
## writing ... [0 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

- (m) Change the values of minimum support and minimum confidence and see how does that affect the number of rules generated.

```
rules <- apriori(Groceries, parameter = list(supp = 0.01, conf = 0.5))

## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##      0.5      0.1      1 none FALSE          TRUE      5      0.01      1
## maxlen target  ext
##      10 rules FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE  FALSE TRUE      2      TRUE
##
## Absolute minimum support count: 98
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.01s].
## sorting and recoding items ... [88 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 done [0.00s].
## writing ... [15 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

- (n) Obtain the association rules with minsup=0.01 and minconf=0.25. Using the functions summary, quality, plot and inspect acquire more information on the generated rules.

```
rules <- apriori(Groceries, parameter = list(supp = 0.01, conf = 0.25))

## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##      0.25      0.1      1 none FALSE          TRUE      5      0.01      1
## maxlen target  ext
##      10 rules FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE  FALSE TRUE      2      TRUE
##
## Absolute minimum support count: 98
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[169 item(s), 9835 transaction(s)] done [0.01s].
## sorting and recoding items ... [88 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 done [0.00s].
## writing ... [171 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].

summary(rules)

## set of 171 rules
##
## rule length distribution (lhs + rhs):sizes
##  1  2  3
##  1 96 74
##
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  1.000  2.000  2.000  2.427  3.000  3.000
##
## summary of quality measures:
```

```
##      support      confidence      lift      count
## Min. :0.01007 Min. :0.2517 Min. :0.9932 Min. : 99.0
## 1st Qu.:0.01159 1st Qu.:0.2965 1st Qu.:1.5175 1st Qu.: 114.0
## Median :0.01454 Median :0.3582 Median :1.7716 Median : 143.0
## Mean :0.01961 Mean :0.3697 Mean :1.8695 Mean : 192.9
## 3rd Qu.:0.02115 3rd Qu.:0.4252 3rd Qu.:2.1412 3rd Qu.: 208.0
## Max. :0.25552 Max. :0.5862 Max. :3.2950 Max. :2513.0
##
## mining info:
##      data ntransactions support confidence
## Groceries      9835      0.01      0.25

inspect(rules[1:5])

##      lhs      rhs      support      confidence lift
## [1] {} => {whole milk} 0.25551601 0.2555160 1.000000
## [2] {hard cheese} => {whole milk} 0.01006609 0.4107884 1.607682
## [3] {butter milk} => {other vegetables} 0.01037112 0.3709091 1.916916
## [4] {butter milk} => {whole milk} 0.01159126 0.4145455 1.622385
## [5] {ham} => {whole milk} 0.01148958 0.4414062 1.727509
##      count
## [1] 2513
## [2] 99
## [3] 102
## [4] 114
## [5] 113
```

- (o) Select the rules with a lift value above 2. Use the function subset for that.

```
rules.sub <- subset(rules, subset = lift > 2)
inspect(rules.sub[1:5])

##      lhs      rhs      support      confidence lift
## [1] {onions} => {other vegetables} 0.01423488 0.4590164 2.372268
## [2] {berries} => {yogurt} 0.01057448 0.3180428 2.279848
## [3] {hamburger meat} => {other vegetables} 0.01382816 0.4159021 2.149447
## [4] {cream cheese } => {yogurt} 0.01240468 0.3128205 2.242412
## [5] {chicken} => {root vegetables} 0.01087951 0.2535545 2.326221
##      count
## [1] 140
## [2] 104
## [3] 136
## [4] 122
## [5] 107
```


- (p) Using one instruction only, select the rules that have lift value above 2 and the items "whole milk" or "yogurt" on the consequent. Inspect the selected rules by decreasing order of their lift value.

```
rules.sub <- subset(rules, subset = rhs %in% c("yogurt", "whole milk") & lift >
2)
rules.sort <- sort(rules.sub, by = "lift")
inspect(rules.sort[1:5])

##      lhs      rhs      support
## [1] {whole milk,curd} => {yogurt} 0.01006609
## [2] {tropical fruit,whole milk} => {yogurt} 0.01514997
## [3] {other vegetables,whipped/sour cream} => {yogurt} 0.01016777
## [4] {tropical fruit,other vegetables} => {yogurt} 0.01230300
## [5] {whole milk,whipped/sour cream} => {yogurt} 0.01087951
##      confidence lift      count
## [1] 0.3852140 2.761356 99
## [2] 0.3581731 2.567516 149
## [3] 0.3521127 2.524073 100
## [4] 0.3427762 2.457146 121
## [5] 0.3375394 2.419607 107

plot(rules.sub)
```


Scatter plot for 23 rules



```
plot(rules.sub, method = "graph")
```

Graph for 23 rules

size: support (0.01 – 0.029)
color: lift (2 – 2.761)



2. Read the csv file of **German Credit dataset** into a data frame in R. This data set has the record of 1000 persons who took a credit by a bank.

```
df <- tbl_df(read.csv("german_credit.csv"))
```

(a) Remove the first attribute from the data frame, it is just an identifier for each record.

```
df <- df %>% select(-default)
```

(b) Try to convert the data frame into a transactions data set using the function `as`. What do you obtain?

```
dfT <- as(df, "transactions")

## Error in discretizeDF(from): Problem with column installment.as.income.perc
## Error in discretize(x = c(4L, 2L, 2L, 2L, 3L, 2L, 3L, 2L, 2L, 4L, 3L, :
## The calculated breaks are: 1, 2, 4, 4
## Some breaks are not unique. Change the number of breaks or consider using method 'fixed'.
```

(c) Use the function `cut` to discretize the numerical attributes according to the following:

- `duration_in_month`: 4 equal-with intervals with labels "short", "med-short", "med-long", "long";
- `credit_amount`: 4 equal-with intervals with labels "small", "med-small", "med-high", "high";
- `age`: 4 equal-with intervals with labels "young adult", "adult", "senior", "golden".
- to the rest of numerical attributes, simply use the function `as.factor`

```
df <- df %>% mutate(duration_in_month = cut(duration_in_month, 4, labels = c("short",
"med-short", "med-long", "long")), credit_amount = cut(credit_amount, 4,
labels = c("small", "med-small", "med-high", "high")), age = cut(age, 4,
labels = c("young adult", "adult", "senior", "golden")))

df <- df %>% mutate_if(is.numeric, as.factor)
```

(d) Convert the data frame into a data set of transactions. What do you obtain? Use the function `itemInfo` to see what each item represents.

```
dfT <- as(df, "transactions")
item_dfT <- itemInfo(dfT)
head(item_dfT)

##                                labels
## 1                                account_check_status=< 0 DM
## 2 account_check_status>= 200 DM / salary assignments for at least 1 year
## 3                                account_check_status=0 <= ... < 200 DM
## 4                                account_check_status=no checking account
## 5                                duration_in_month=short
## 6                                duration_in_month=med-short
##                                variables                                levels
## 1 account_check_status                                < 0 DM
## 2 account_check_status >= 200 DM / salary assignments for at least 1 year
## 3 account_check_status                                0 <= ... < 200 DM
## 4 account_check_status                                no checking account
## 5 duration_in_month                                short
## 6 duration_in_month                                med-short
```

(e) Run `apriori` to obtain the association rules from the data set. Plot the obtained rules.


```
rules <- apriori(dfT)

## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
```



```
##      0.8    0.1    1 none FALSE      TRUE    5    0.1    1
## maxlen target ext
##      10 rules FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 100
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[80 item(s), 1000 transaction(s)] done [0.00s].
## sorting and recoding items ... [53 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 5 6 7 8 9 10 done [0.06s].
## writing ... [46642 rule(s)] done [0.01s].
## creating S4 object ... done [0.02s].

plot(rules)
```



(f) Observe the effect of filters and measures on the number of rules generated.

(g) Select the rules with confidence equal to 1. What does those rules tell you?

```
rules.conf1 <- subset(rules, confidence == 1)
```

(h) Run apriori again, but this time imposing a minimum confidence equal to 0.6, minimum length of 2 and focusing only on attributes sex, age, job, housing and purpose of credit.

```
myItems <- subset(item_dfT, variables %in% c("age", "personal_status_sex", "job", "housing", "purpose"))$labels
rules <- apriori(dfT,
  parameter = list(conf=0.6, minlen=2), # 44 rules
  appearance = list(both = myItems,
    default="none"))

## Apriori
##
## Parameter specification:
## confidence minval smax arem avar originalSupport maxtime support minlen
##      0.6    0.1    1 none FALSE      TRUE    5    0.1    2
## maxlen target ext
##      10 rules FALSE
##
```

```
## Algorithmic control:
## filter tree heap memopt load sort verbose
## 0.1 TRUE TRUE FALSE TRUE 2 TRUE
##
## Absolute minimum support count: 100
##
## set item appearances ...[25 item(s)] done [0.00s].
## set transactions ...[25 item(s), 1000 transaction(s)] done [0.00s].
## sorting and recoding items ... [15 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 done [0.00s].
## writing ... [44 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
```

- (i) Identify rules $a \rightarrow b$ and $b \rightarrow a$. What do their quality values tell you?

```
# same lift and same support different confidence [19] {job=skilled employee
# / official} => {housing=own} 0.452 0.7174603 1.0062557 [20] {housing=own}
# => {job=skilled employee / official} 0.452 0.6339411 1.0062557 housing =
# own appears more often in transactions that contain job = skilled
```

- (j) Run apriori to obtain rules that relate the purpose of credit with age, job and housing. Impose a minimum support of 0.05, minimum confidence of 0.25 and a minimum length of 2. Could you propose a marketing campaign from the obtained rules?

```
my.lhs <- subset(item_dft, variables %in% c("age", "job", "housing"))$labels
my.rhs <- subset(item_dft, variables == "purpose")$labels
rules1 <- apriori(dft, parameter = list(confidence = 0.25, minlen = 2, support = 0.05),
  appearance = list(lhs = my.lhs, rhs = my.rhs, default = "none"))

## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
## 0.25 0.1 1 none FALSE TRUE 5 0.05 2
## maxlen target ext
## 10 rules FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
## 0.1 TRUE TRUE FALSE TRUE 2 TRUE
##
## Absolute minimum support count: 50
##
## set item appearances ...[21 item(s)] done [0.00s].
## set transactions ...[21 item(s), 1000 transaction(s)] done [0.00s].
## sorting and recoding items ... [15 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 done [0.00s].
## writing ... [13 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].

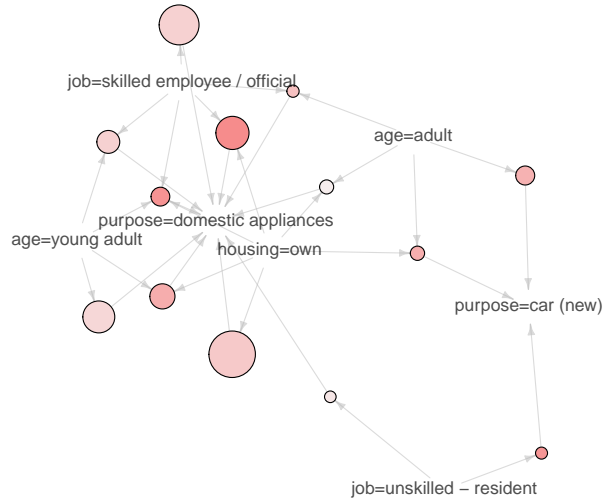
# promote credit for domestic appliances among young adults or adults with
# skilled job and own housing
```

- (k) Plot the previous set of rules using the method graph and graph with itemsets. What do these graphs tell you?

```
plot(rules1, method = "graph")
```

Graph for 13 rules

size: support (0.057 – 0.227)
color: lift (0.99 – 1.28)

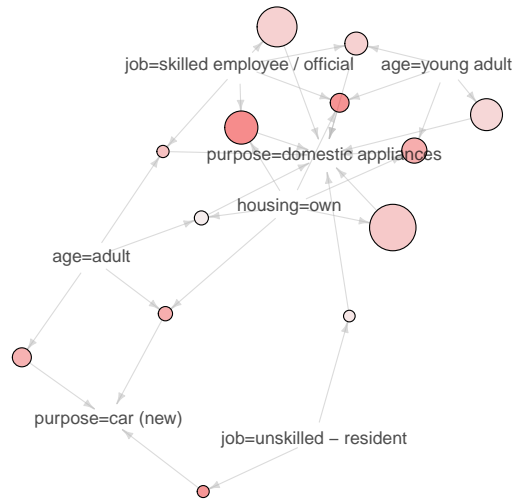


```
plot(rules1, method = "graph", control = list(type = "itemsets"))

## Available control parameters (with default values):
## main = Graph for 13 rules
## nodeColors = c("#66CC6680", "#9999CC80")
## nodeCol = c("#EE0000FF", "#EE0303FF", "#EE0606FF", "#EE0909FF", "#EE0C0CFF", "#EE0F0FFF", "#EE1212FF", "#EE1515FF", "#EE1818FF",
## edgeCol = c("#474747FF", "#494949FF", "#4B4B4BFF", "#4D4D4DFF", "#4F4F4FFF", "#515151FF", "#535353FF", "#555555FF", "#575757FF",
## alpha = 0.5
## cex = 1
## itemLabels = TRUE
## labelCol = #000000B3
## measureLabels = FALSE
## precision = 3
## layout = NULL
## layoutParams = list()
## arrowSize = 0.5
## engine = igraph
## plot = TRUE
## plot_options = list()
## max = 100
## verbose = FALSE
```

Graph for 13 rules

size: support (0.057 – 0.227)
color: lift (0.99 – 1.28)



(I) Plot the previous set of rules using the method grouped.

```
plot(rules1, method = "grouped")
```

Grouped Matrix for 13 Rules

